

Bandit on the loose!

Networking in Elixir demystified

@mattrudel - github.com/mtrudel - mat@geeky.net

github.com/mtrudel/talks

Introducing Thousand Island & Bandit

github.com/mtrudel/thousand_island - github.com/mtrudel/bandit

Phoenix
Plug
HTTP
Sockets
TCP/IP

Phoenix
Plug
HTTP
Sockets
TCP/IP

In the beginning there was IP (RFC 791)

- Provides each host on the network with an IP address
- 'Being connected to the internet' means¹ having an IP yourself and being able to reach any other host via their IP
- All IP does is deliver packets to a host
- IP 'encapsulates' higher level protocols such as TCP and UDP
 - Identified by a unique 'Protocol Number' (TCP is 6, UDP is 17)

¹A million caveats to this (NAT, the DFZ, *cast, etc)

Then came TCP (RFC 793)

- Provides application-level addressability
- Addressed by an IP & a port number
- Connections come in the form of a pair of reliable streams
- Each direction is a one-way stream of binary data
- Each direction is independent, can be shutdown, used simultaneously
- Commonly accessed via the 'Berkeley Socket API', part of libc

Phoenix
Plug
HTTP
Sockets
TCP/IP

Phoenix

Plug

HTTP

Sockets

TCP/IP

Berkeley sockets in a nutshell

- Clients **connect** to an IP/port pair
- Servers **bind** to a named port, and **listen** for connections
- Servers then **accept** clients' **connect** requests in a loop
- Once accepted, a TCP session is created between the client and the server
- The two peers can now **send** and **recv** data with one another
- At any time, either peer can **close** a connection

How does Elixir access sockets?

- Via Erlang's `:gen_tcp` (and `:socket` since OTP 22)
- Provides nearly 1:1 mapping onto the standard Berkeley calls
- SSL is accessed via `:ssl` using ~identical functions
 - All the core socket abstractions are the same as TCP
 - `s/:gen_tcp/:ssl/`

Simple :gen_tcp Server


```
{:ok, listen_socket} = :gen_tcp.listen(4000, [active: false])  
accept_and_handle(listen_socket)
```

```
def accept_and_handle(listen_socket) do  
  {:ok, socket} = :gen_tcp.accept(listen_socket)  
  :gen_tcp.send(socket, "Hello, World")  
  :gen_tcp.close(socket)  
  accept_and_handle(listen_socket)  
end
```

Simple :gen_tcp Client

```
{:ok, socket} = :gen_tcp.connect('localhost', 4000, [active: false])
{:ok, data} = :gen_tcp.recv(socket, 0)
IO.puts(data) #=> "Hello, World"
:gen_tcp.close(socket)
```

The Stack So Far



Socket Review

- This is all you **really** need to implement a server
- Working this low level is tedious
- Performant patterns are non-trivial
- Usually abstracted by a 'socket server'

What is a Socket Server?

- Provides a useful abstraction of raw sockets
- Sits underneath an e.g. HTTP server
- Listens for client connections over TCP/SSL
- Hands individual connections off to an upper protocol layer
- Handles transport concerns (SSL/TLS negotiation, connection draining, etc)
- Does this efficiently and scalably

Who is a Socket Server?

- Ranch is the current goto implementation of such a server on the BEAM
- Thousand Island² is a new pure-Elixir socket server, inspired by Ranch but with numerous improvements

²Get it? They're both salad dressings?!?

Thousand Island

A pure Elixir socket server, embodying OTP ideals

- Fast (performance basically equal to Ranch)
- Simple (1.7k LoC, about ½ the size of Ranch)
- Comprehensive documentation & examples
- Supports TCP, TLS & Unix Domain sockets
- Loads of config options
- Fully wired for telemetry, including socket-level tracing
- Extremely simple & powerful Handler behaviour

Let's build RFC867 (Daytime)

```
defmodule Daytime do
  use ThousandIsland.Handler

  @impl ThousandIsland.Handler
  def handle_connection(socket, state) do
    time = DateTime.utc_now() |> to_string()
    ThousandIsland.Socket.send(socket, time)
    {:close, state}
  end
end

{:ok, pid} = ThousandIsland.start_link(handler_module: Daytime)
```

Let's build RFC862 (Echo)

```
defmodule Echo do
  use ThousandIsland.Handler

  @impl ThousandIsland.Handler
  def handle_data(data, socket, state) do
    ThousandIsland.Socket.send(socket, data)
    {:continue, state}
  end
end
```

```
{:ok, pid} = ThousandIsland.start_link(handler_module: Echo)
```

Thousand Island Handler Process

- 1 process per connection
- **ThousandIsland.Handler** module is just a **GenServer**
- Your connection processes can do all the normal **GenServer** things
- Receives are async by default (via **handle_data/3**)
- Free to write 'traditional' blocking network code if you wish, anywhere
- **start_link** based escape hatch if you need it (you probably won't)

Thousand Island Process Model

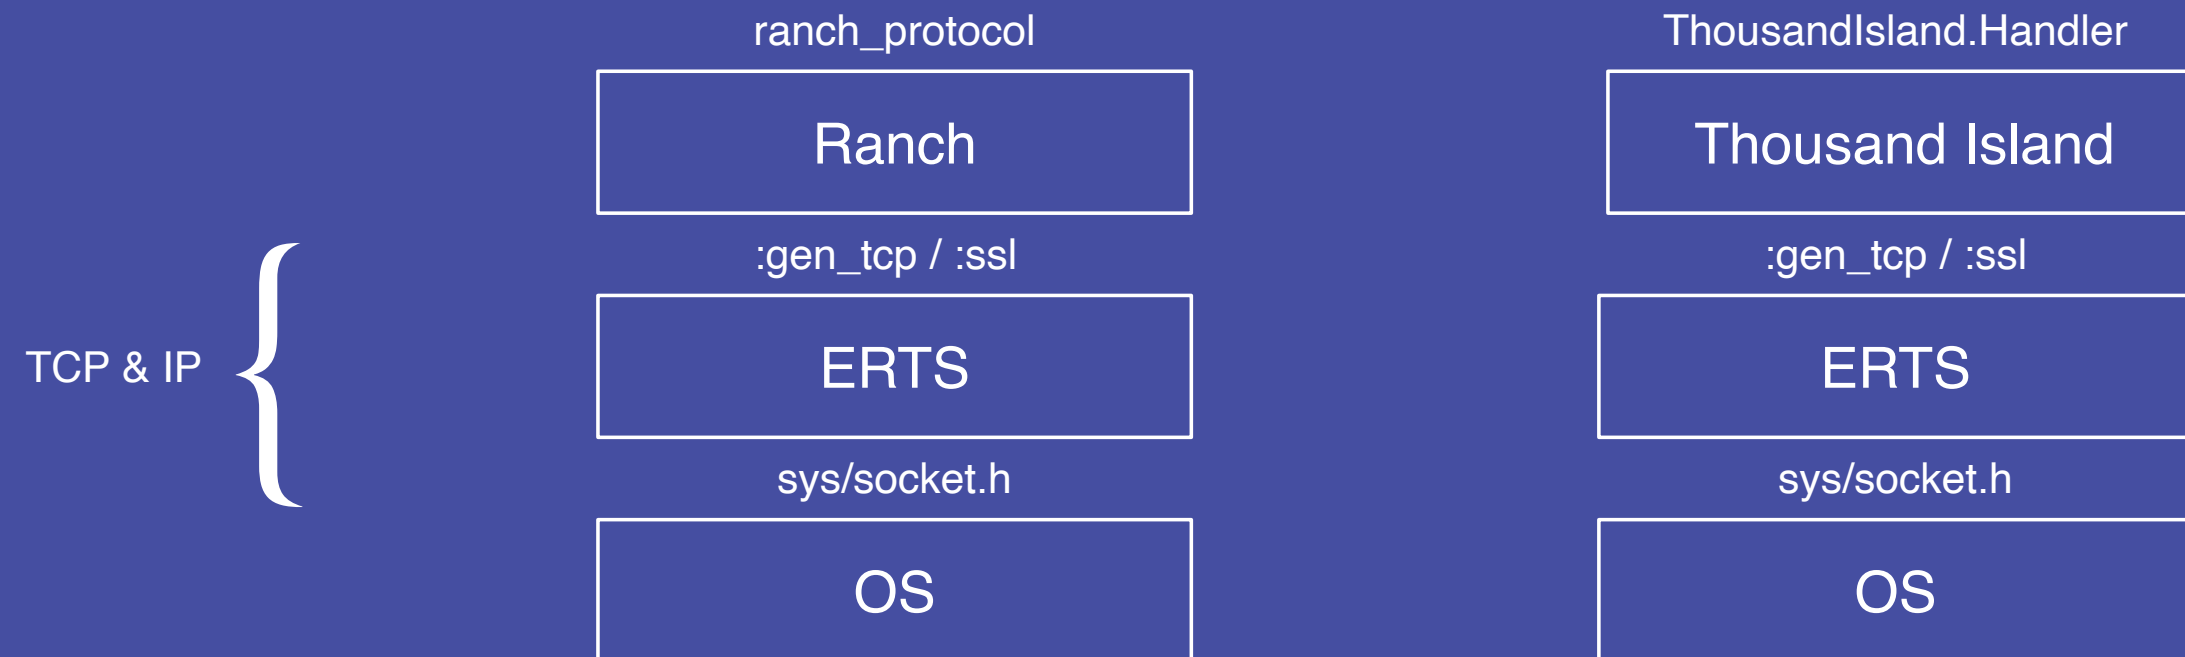
- Handler processes are hosted within a process tree
- Rooted in the **ThousandIsland.start_link/1** call which starts the server
- Process tree is entirely self-contained
- Multi-level process tree designed to minimize contention
- Textbook example of how powerful OTP design patterns can be
- Check out the project README for more info

Thousand Island

github.com/mtrudel/thousand_island

- Stable and suitable for general use
- 0.5.x series is likely the last before a 1.0
- Big ideas for the future, but the foundation is solid

The Stack So Far



Phoenix

Plug

HTTP

Sockets

TCP/IP

What is HTTP?

- HTTP/1.x is just plain text over sockets. Defined primarily by RFC 2616

What is HTTP?

- HTTP/1.x is plain text over sockets. Defined primarily by RFC 2616

```
> GET /thing/12 HTTP/1.1  
> host: www.example.com  
> [... other headers ...]  
> [empty line]  
> [body content]
```

```
< HTTP/1.1 200 OK  
< content-length: 123  
< [... other headers ...]  
< [empty line]  
< [body content]
```

Let's build a simple HTTP server

```
defmodule HelloWorldHTTP do
  use ThousandIsland.Handler

  @impl ThousandIsland.Handler
  def handle_data(_data, socket, state) do
    ThousandIsland.Socket.send(socket, "HTTP/1.0 200 OK\r\n\r\nHello, World")
    {:close, state}
  end
end
```

What does an HTTP server do?

- A complete HTTP/1.x implementation isn't **that** much more complicated
- Parses & validates request and header lines (and possibly request body)
- Sends conformant responses back to the client

Who is an HTTP server?

- Cowboy is the current goto HTTP server on the BEAM
 - Complete HTTP/1.1, HTTP/2 & WebSocket server
- I've written Bandit³, a new pure-Elixir HTTP server for Plug applications

³Get it? They're both tropes within a heavily romanticized & largely fictional American mythology?!?

Bandit

An HTTP Server for Plug Applications

- Written 100% in Elixir
- Plug-native
- Robust HTTP/1.1 and HTTP/2 conformance
- Written from the ground up for correctness, performance & clarity
- Incredible performance (more later!)

github.com/mtrudel/bandit

Making an HTTP server do useful things

- So, HTTP servers return content to clients, but **what** content?
- The content from our app, obviously
- But how to hand off to app logic?
- Enter Plug

Phoenix

Plug

HTTP

Sockets

TCP/IP

Plug

An abstraction of HTTP Request/Response Pairs

```
defmodule HelloWorldPlug do
  def init(opts), do: opts

  def call(%Plug.Conn{} = conn, _opts) do
    Plug.Conn.send_resp(conn, 200, "Hello, World")
  end
end
```

Plug Comes With Batteries Included

- Plug comes with built-in Plugs to:
 - Create pipelines
 - Route requests
 - Manage sessions
 - Parse request bodies
 - **Lots** more
- This is enough for a LOT of applications

Aside: Bandit & Thousand Island's Origin Story

- Written to support HAP, a HomeKit Accessory Protocol library for Nerves
- Runs over HTTP, but with a twist (custom encryption on bare TCP)

Aside: Bandit & Thousand Island's Origin Story

- Written to support HAP, a HomeKit Accessory Protocol library for Nerves
- Runs over HTTP, but with a twist (custom encryption on bare TCP)
- HTTP layer is simple, runs entirely on **Plug.Router**:

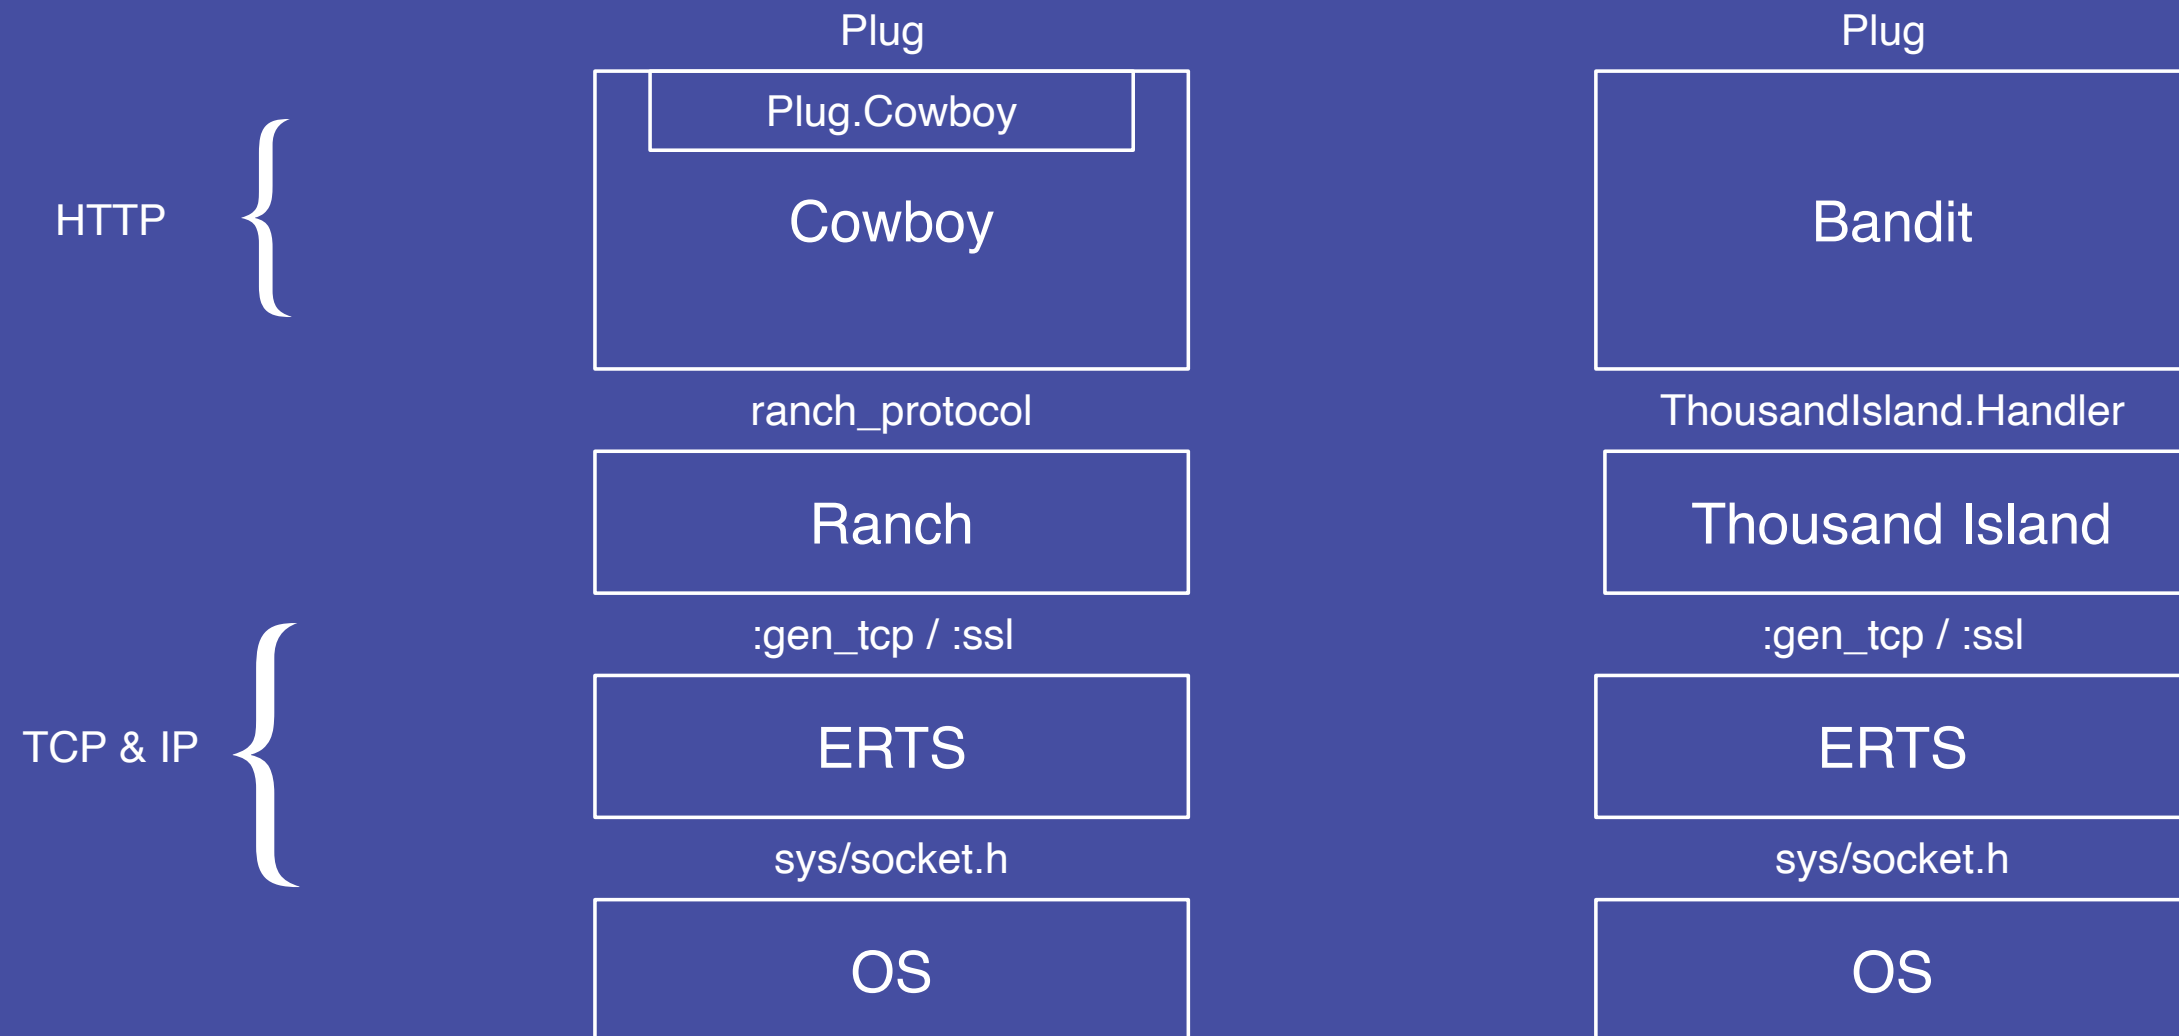
```
defmodule HAP.EncryptedHTTPServer do
  use Plug.Router

  post "/pairings" do ....
  get  "/accessories" do ....
end
```

Plug Servers

- Plug is implemented in Cowboy via the **Plug.Cowboy** adapter
- Bandit is Plug-Native

The Stack So Far



Phoenix

Plug

HTTP

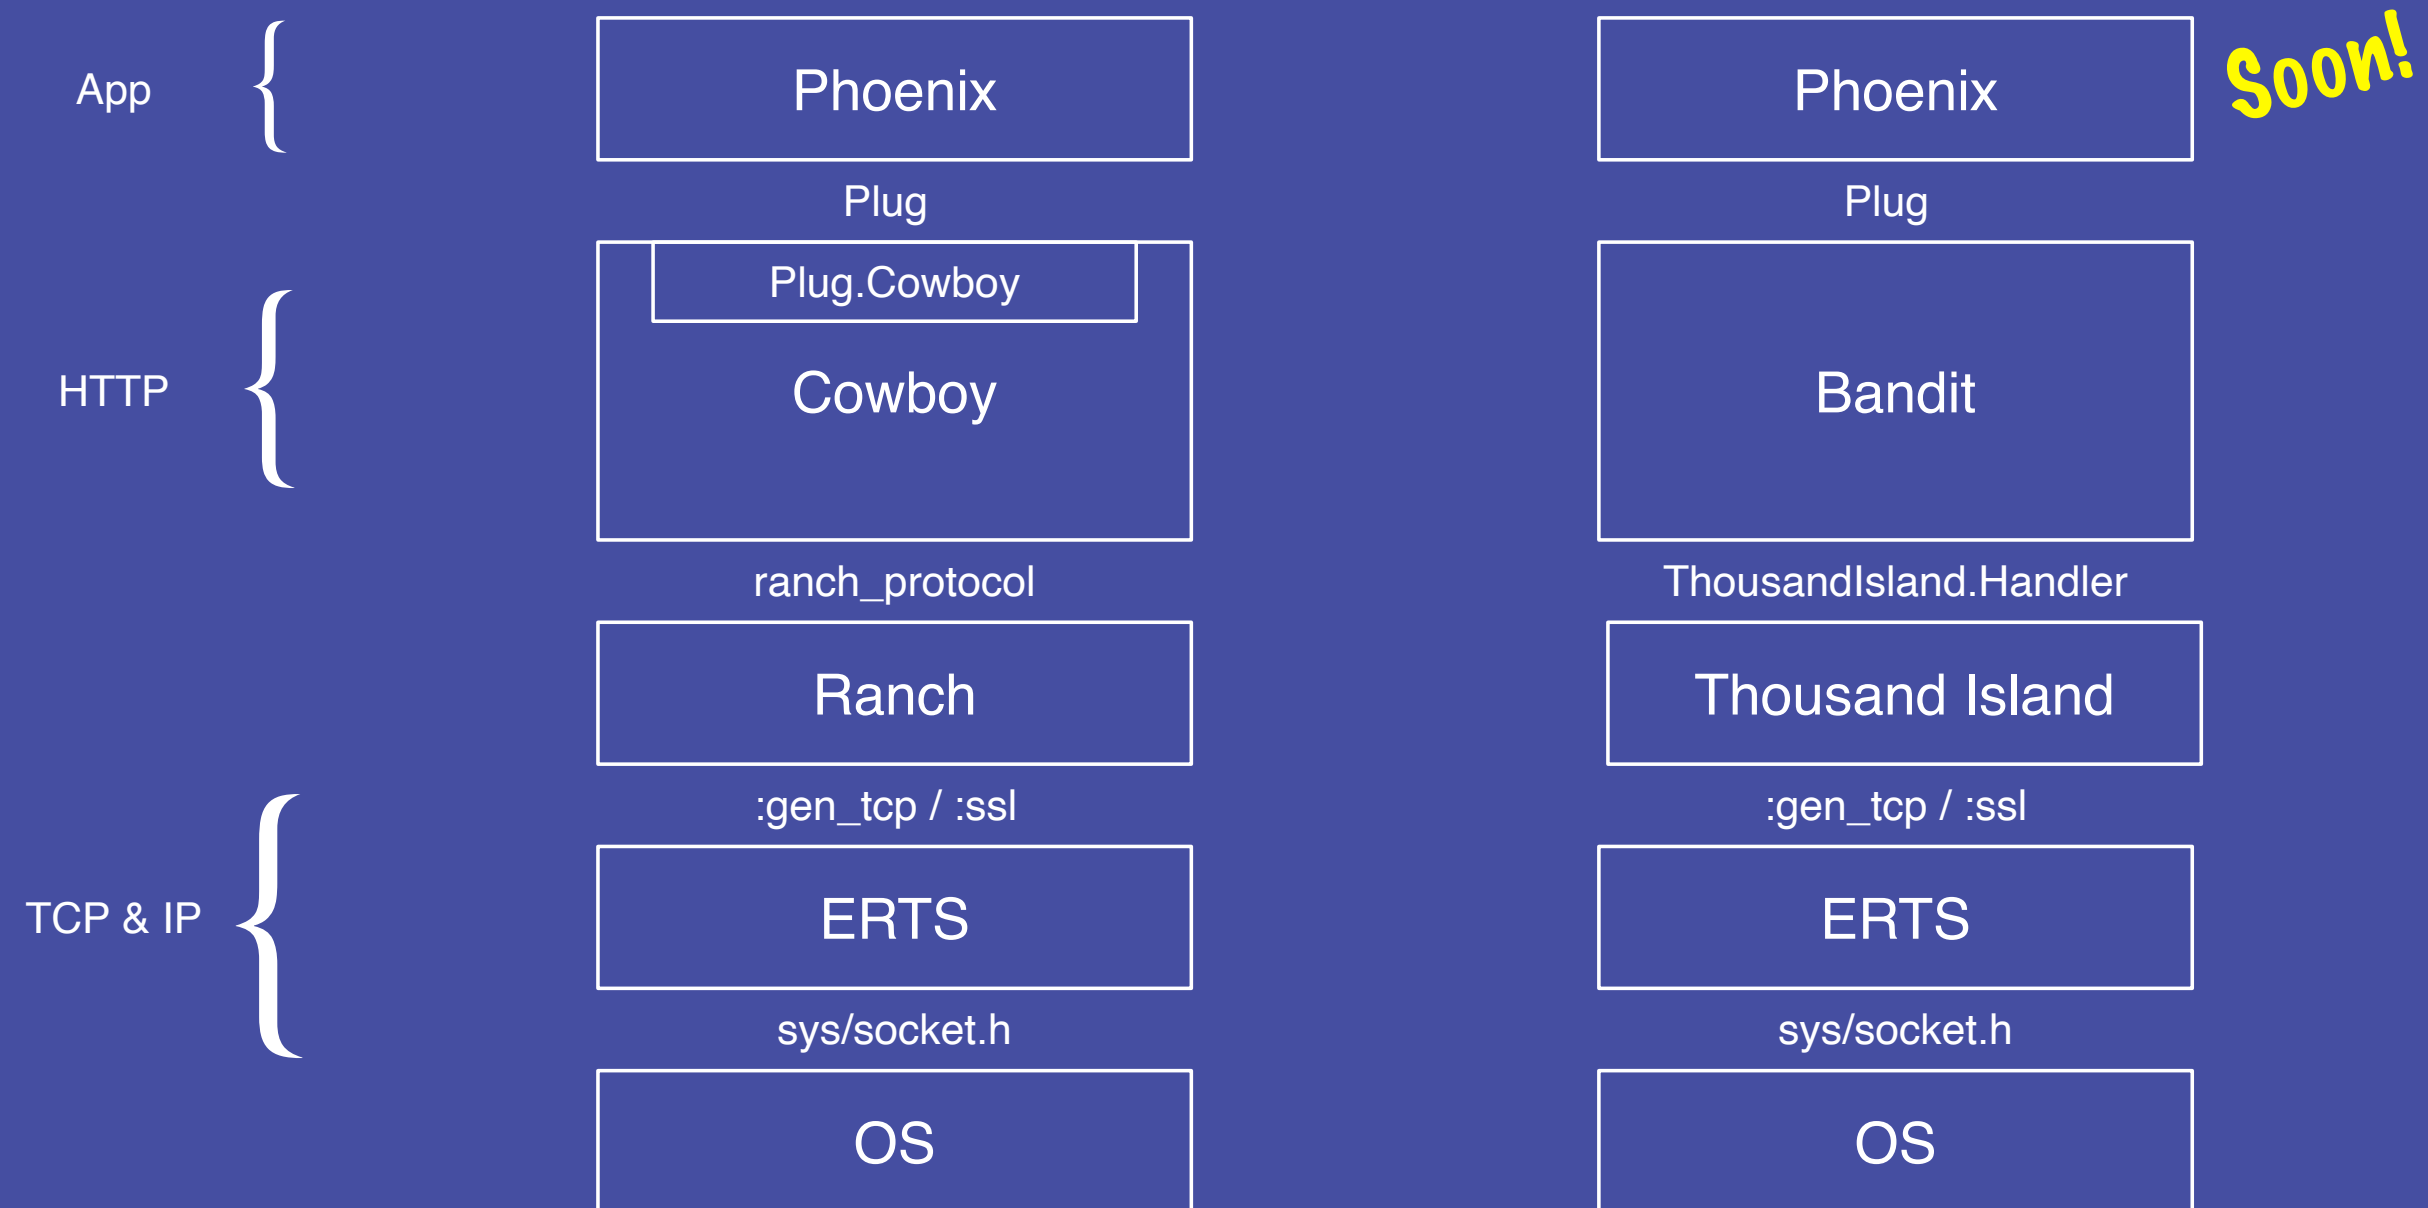
Sockets

TCP/IP

Enter Phoenix

- The HTTP part of Phoenix is 'just' a Plug
- Phoenix's WebSocket support is a different thing (more later)
- Not much more to say about things here
- This is as high in the stack as we're going today

The Full Stack (Finally)



Bandit

An HTTP Server for Plug Applications

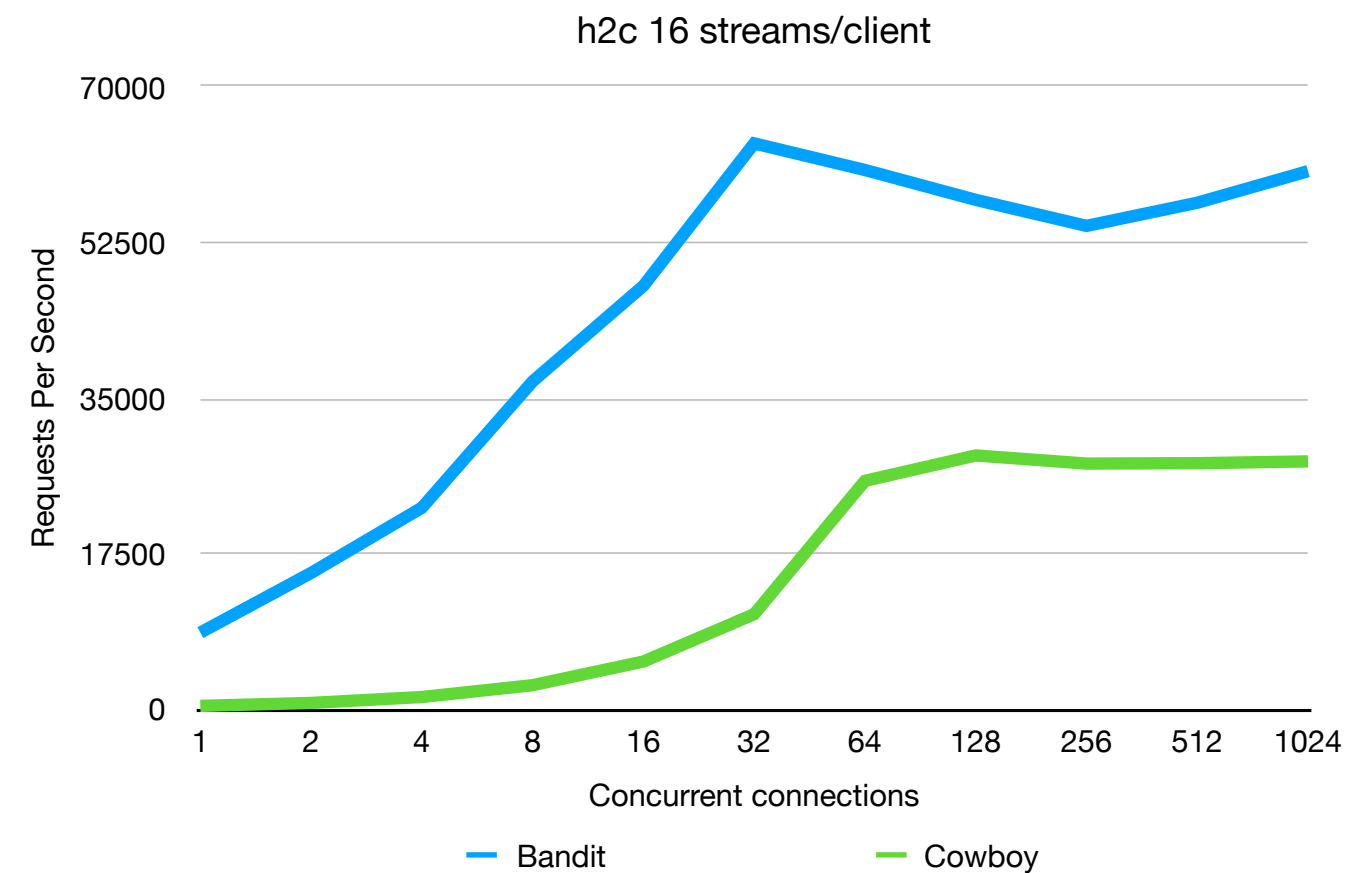
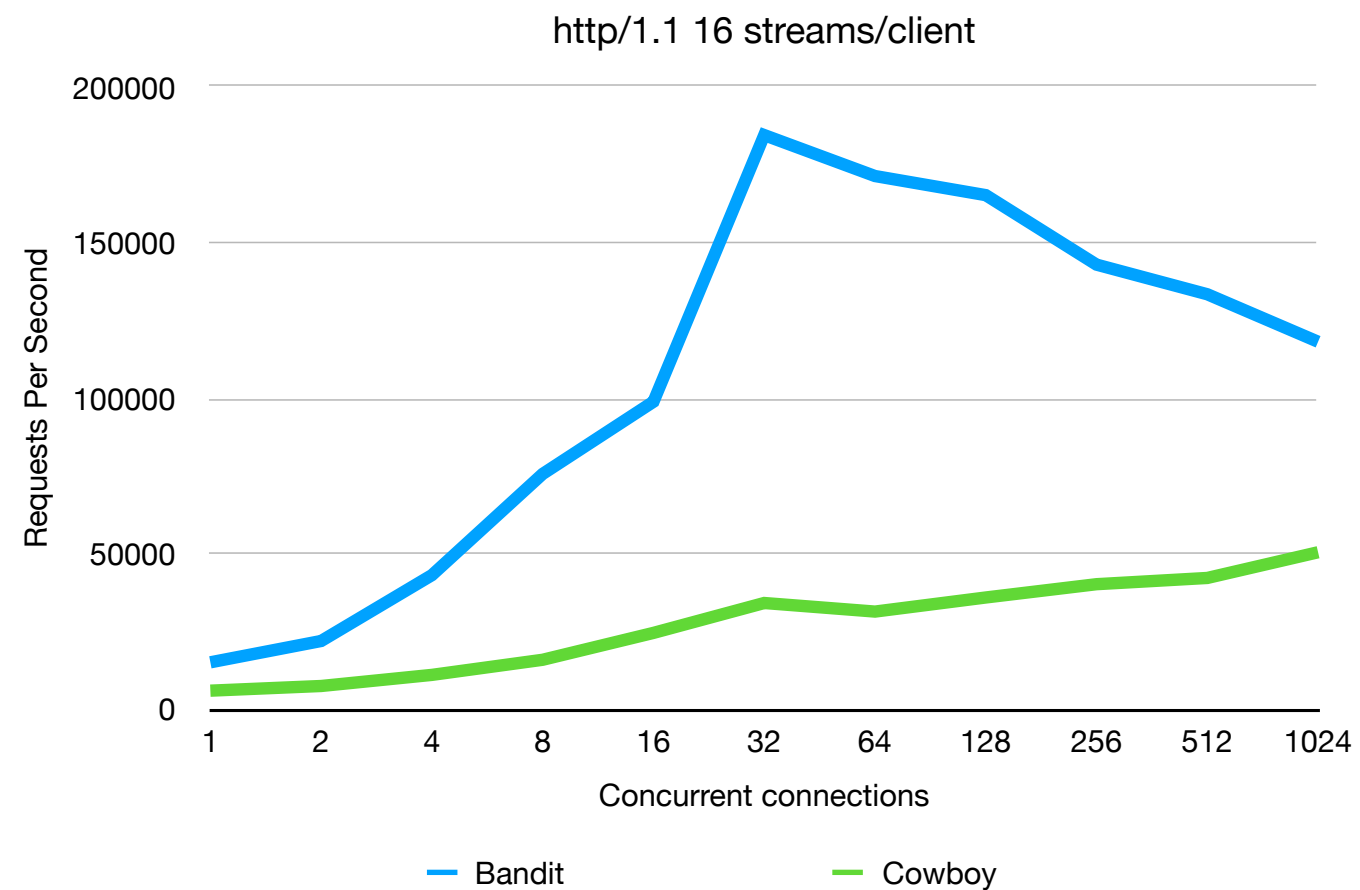
- Built on top of Thousand Island
- Full support for HTTP/1.x & HTTP/2
 - Scores 100% on h2spec in **--strict** mode (runs in CI)
- Clear, Approachable, Idiomatic Elixir

github.com/mtrudel/bandit

Bandit is a Plug-First Server

- No impedance mismatches
- Less (and clearer) code (<3k LoC today, about 1/3 of Cowboy)
- **Incredibly** fast and memory efficient

Up to 5x Faster Than Cowboy



HTTP/1.1 In Bandit

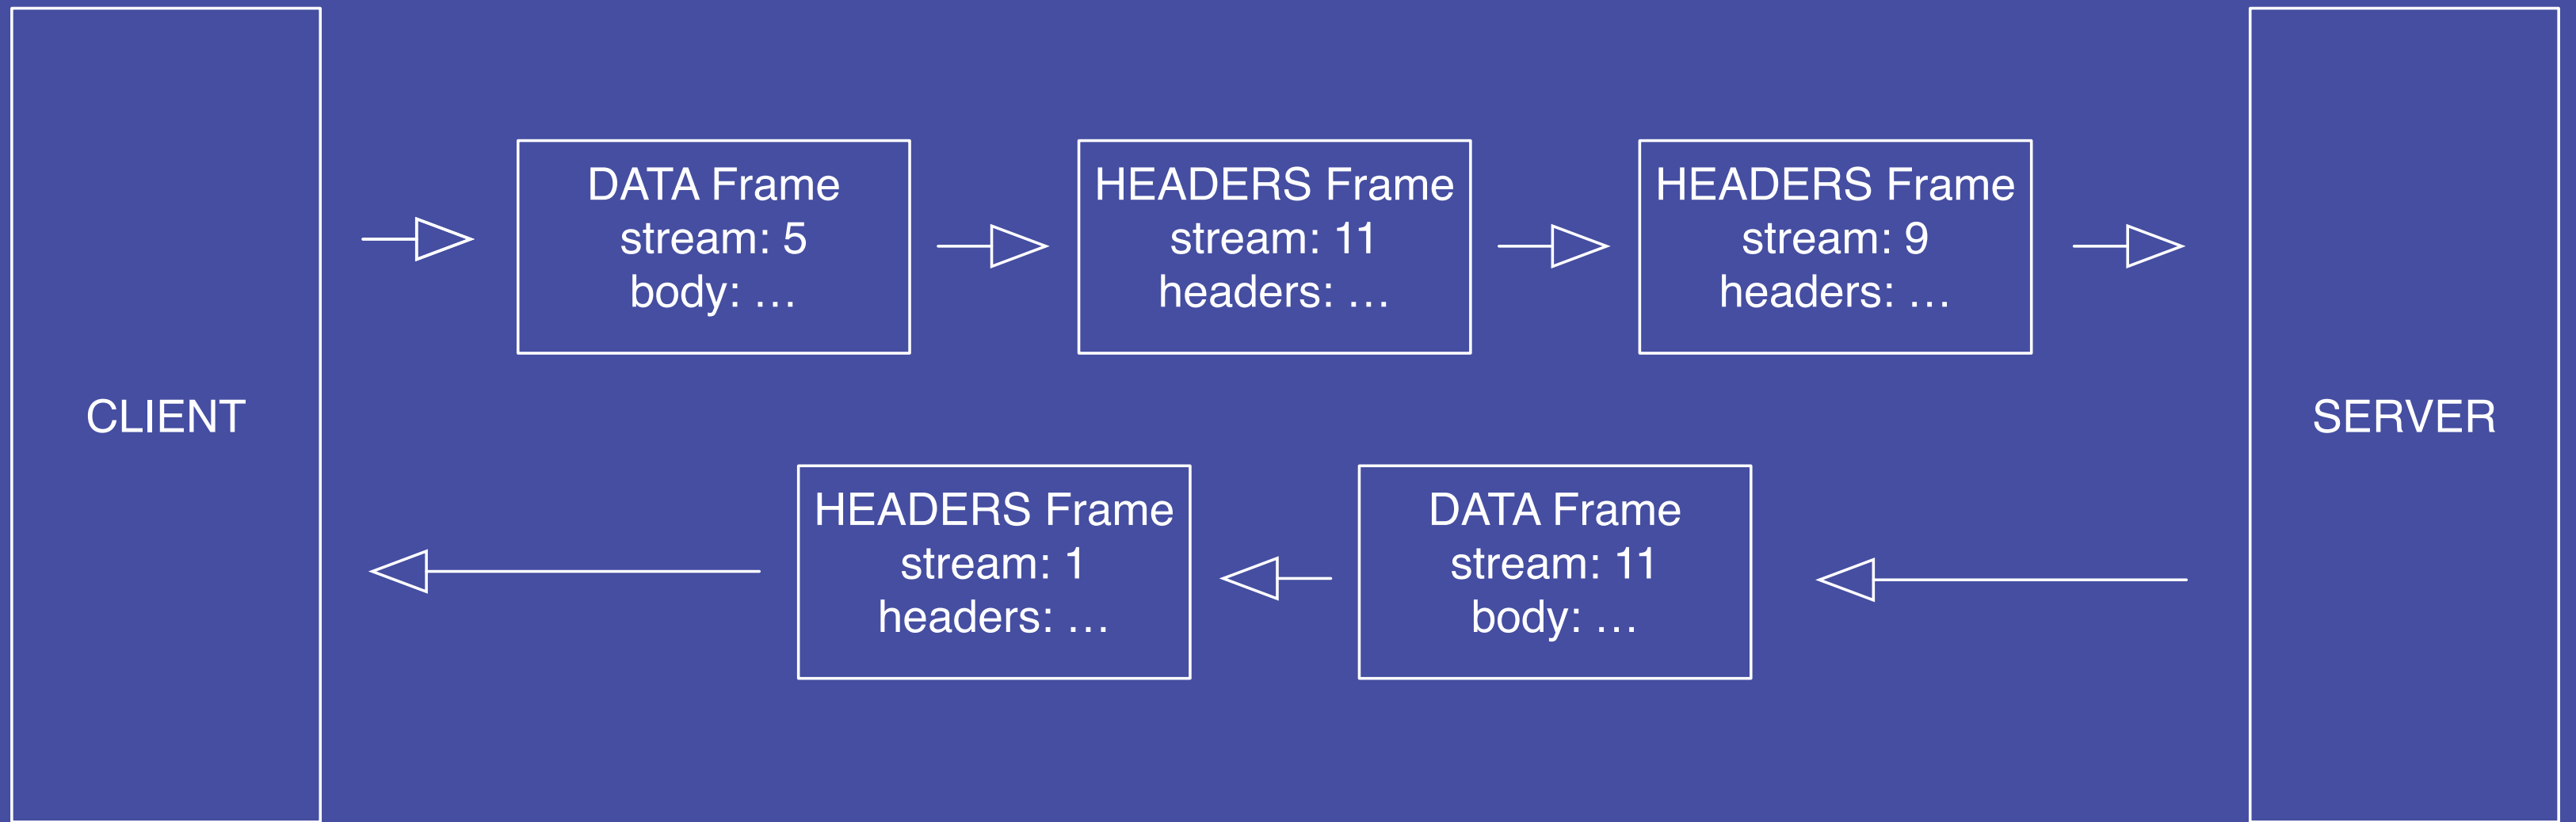
- 1 process per connection
- 1:1 mapping between Handler `handle_data/3` calls and Plug `call/2` calls
- Straightforward, 'linear' code

HTTP/2 In Bandit

- Quite a bit harder

HTTP/2 In A Nutshell

Framed binary protocol



HTTP/2 In Bandit

- 1 process per connection + 1 process per stream
- Connection process implements **ThousandIsland.Handler**
- Stream processes are the ones that make Plug calls
- Full support for flow control, push promises, all the goodies
- `iodata` everywhere for speed and memory awesomeness

Where Is Bandit Today?

- HTTP/2 implementation complete & exhaustively tested (0.3.x)
- HTTP/1.1 implementation undergoing refactor & test improvements (0.4.x)
- Suitable for non-production Plug apps (not Phoenix, yet)
- Drop-in replacement for Cowboy where appropriate

Future Plans for Bandit

- Phoenix Support
 - Requires WebSocket support (coming in 0.5.x)
 - Note that WebSocket support is outside the scope of Plug
 - Requires Phoenix integration (coming in 0.7.x)
- Continue to improve perf and Cowboy migration story
- 'PR Campaign' to drive adoption
- HTTP/3 Support (distant future)

The Goal:

The Goal:

**Become the de-facto
networking stack
for Elixir & the BEAM
(along with Thousand Island)**

Working With Our Friends

- Work within Phoenix needed to add Bandit support
 - Possibly generalize its WebSocket interface
- Code sharing with the Mint HTTP client (Thanks Andrea & Eric for HPAX)
- Improve iolist primitives in ERTS
- Other minor upstream improvements (eg. URI canonicalization & trailer support in Plug)
- To be clear: Cowboy & Ranch are awesome projects

Contribute!

- Fun, foundational work
- Grassroots project
- Contributions are **extremely** welcome
 - Core HTTP/1.1 & WebSocket work
 - Property-based testing suite
 - Security review & mitigation
 - Profiling & perf improvements / automated testing
 - **Lots** more!

github.com...

/mtrudel/thousand_island

/mtrudel/bandit

/mtrudel/talks

@mattrudel (Twitter)

@mtrudel (Elixir Slack)

mat@geeky.net