

# Enstabilize

&

# Quickify

Yer Tests

# To Do Better Stuff

Mat Trudel // [mat@geeky.net](mailto:mat@geeky.net) // [@mtrudel](https://twitter.com/mtrudel) // [github.com/mtrudel/talks](https://github.com/mtrudel/talks)

**Tests are the  
foundation  
of your work**

(and you should strive to have good ones)

**What makes  
for  
good tests?**

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable**

Tests should be **ubiquitous**

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable**

Tests should be **ubiquitous**

# Tests should be clear

Consistent Setup / Trigger / Assertion structure

```
test "getting user details" do
  # Setup
  user = user_fixture()

  # Trigger
  response = Req.get("/users/#{user.id}")

  # Assertions
  assert response.name == user.name
end
```

# Tests should be clear

## Consistent Setup / Trigger / Assertion structure

```
test "getting user details" do
  # Setup
  user = user_fixture()

  # Trigger
  response = Req.get("/users/#{user.id}")

  # Assertions
  assert response.name == user.name
end
```

# Tests should be clear

## Consistent Setup / Trigger / Assertion structure

```
test "getting user details" do
  # Setup
  user = user_fixture()

  # Trigger
  response = Req.get("/users/#{user.id}")

  # Assertions
  assert response.name == user.name
end
```



# Tests should be clear

## Consistent Setup / Trigger / Assertion structure

```
test "getting user details" do
  # Setup
  user = user_fixture()

  # Trigger
  response = Req.get("/users/#{user.id}")

  # Assertions
  assert response.name == user.name
end
```

# Tests should be clear

## Factor up repetitive setup

- `ExUnit.Callback` has a ton of useful functionality
  - We all know about `setup` and `setup_all`
  - There's also `start_supervised`
    - Process lifecycle 'just works'
    - 1.17+ `start_supervised` now sets `$callers` and `$ancestors`
  - My personal favourite though, are named setup callbacks
  - A real-life example from Bandit ties all this together:

```
defmodule PlugTest do
  use ServerHelpers

  setup :http_server

  test "the plumbing all works", context do
    response = Req.get!("#{context.base}/hello_world")

    assert response.body == "Hello, World!"
  end

  def call(conn, _opts) do
    send_resp(conn, 200, "Hello, World!")
  end
end

defmodule ServerHelpers do
  defmacro __using__(_) do
    quote do
      import Plug.Conn

      def http_server(_context, opts \\ [plug: __MODULE__]) do
        {:ok, server_pid} = opts |> Bandit.child_spec() |> start_supervised()
        {:ok, {_ip, port}} = ThousandIsland.listener_info(server_pid)
        [base: "http://localhost:#{port}"]
      end
    end
  end
end
```

```
defmodule PlugTest do
  use ServerHelpers

  setup :http_server

  test "the plumbing all works", context do
    response = Req.get!("#{context.base}/hello_world")

    assert response.body == "Hello, World!"
  end

  def call(conn, _opts) do
    send_resp(conn, 200, "Hello, World!")
  end
end

defmodule ServerHelpers do
  defmacro __using__(_) do
    quote do
      import Plug.Conn

      def http_server(_context, opts \\ [plug: __MODULE__]) do
        {:ok, server_pid} = opts |> Bandit.child_spec() |> start_supervised()
        {:ok, {_ip, port}} = ThousandIsland.listener_info(server_pid)
        [base: "http://localhost:#{port}"]
      end
    end
  end
end
```

```
defmodule PlugTest do
  use ServerHelpers

  setup :http_server

  test "the plumbing all works", context do
    response = Req.get!("#{context.base}/hello_world")

    assert response.body == "Hello, World!"
  end

  def call(conn, _opts) do
    send_resp(conn, 200, "Hello, World!")
  end
end

defmodule ServerHelpers do
  defmacro __using__(_) do
    quote do
      import Plug.Conn

      def http_server(_context, opts \\ [plug: __MODULE__]) do
        {:ok, server_pid} = opts |> Bandit.child_spec() |> start_supervised()
        {:ok, {_ip, port}} = ThousandIsland.listener_info(server_pid)
        [base: "http://localhost:#{port}"]
      end
    end
  end
end
```

# Tests should be clear

## Meaningful assertions

- Your assertions should be clear & organized
  - Multiple *related* assertions in a single test are fine
  - Balance between comprehensive & verbose is tricky
- Your tests should be testing what you think they are
  - Beware of the semantic of = vs ==, e.g.:  
    expected = 123  
    assert expected = 123
    - The compiler will *usually* warn
    - ...or you could just use Machete

# Tests should be clear

You should use Machete

- Literate matchers for Elixir
  - Defines the `~>` operator
  - Straightforward literal and var matching
  - Powerful & extensible parameteric matching
  - Robust collection support (including parametric matchers)
  - Useful error messages in ExUnit

# Tests should be clear

You should use Machete

```
assert "abc" ~> "abc"
```

```
assert "abc" ~> string()
```

```
assert 123 ~> integer(positive: true)
```

```
assert "2025-02-12T16:45:00Z" ~> iso8601_datetime(roughly: :now)
```

```
assert %{a: "abc", b: 123} ~> %{a: "abc", b: integer()}
```

```
assert ["def", "abc"] ~> in_any_order(["abc", "def"])
```

```
assert ["abc", "def"] ~> list_of(string(), min: 1, max: 5)
```



```
test "it should send `stop` events for normally completing requests", context do
  Req.get!(context.req, url: "/send_200")

  assert_receive [:telemetry, [:bandit, :request, :stop], measurements, metadata], 500

  assert measurements
    ~> %{
      resp_body_bytes: 0,
      duration: integer(max: System.convert_time_unit(1, :second, :native)),
      monotonic_time: integer(roughly: System.monotonic_time()),
      req_header_end_time: integer(roughly: System.monotonic_time()),
      resp_start_time: integer(roughly: System.monotonic_time()),
      resp_end_time: integer(roughly: System.monotonic_time())
    }

  assert metadata
    ~> %{
      connection_telemetry_span_context: reference(),
      telemetry_span_context: reference(),
      conn: struct_like(Plug.Conn, path_info: ["send_200"]),
      plug: {__MODULE__, []}
    }

end
```

```

test "it should send `stop` events for normally completing requests", context do
  Req.get!(context.req, url: "/send_200")

  assert_receive [:telemetry, [:bandit, :request, :stop], measurements, metadata], 500

  assert measurements
    ~> %{
      resp_body_bytes: 0,
      duration: integer(max: System.convert_time_unit(1, :second, :native)),
      monotonic_time: integer(roughly: System.monotonic_time()),
      req_header_end_time: integer(roughly: System.monotonic_time()),
      resp_start_time: integer(roughly: System.monotonic_time()),
      resp_end_time: integer(roughly: System.monotonic_time())
    }

  assert metadata
    ~> %{
      connection_telemetry_span_context: reference(),
      telemetry_span_context: reference(),
      conn: struct_like(Plug.Conn, path_info: ["send_200"]),
      plug: {__MODULE__, []}
    }

end

```

```

test "it should send `stop` events for normally completing requests", context do
  Req.get!(context.req, url: "/send_200")

  assert_receive [:telemetry, [:bandit, :request, :stop], measurements, metadata], 500

  assert measurements
    ~> %{
      resp_body_bytes: 0,
      duration: integer(max: System.convert_time_unit(1, :second, :native)),
      monotonic_time: integer(roughly: System.monotonic_time()),
      req_header_end_time: integer(roughly: System.monotonic_time()),
      resp_start_time: integer(roughly: System.monotonic_time()),
      resp_end_time: integer(roughly: System.monotonic_time())
    }

  assert metadata
    ~> %{
      connection_telemetry_span_context: reference(),
      telemetry_span_context: reference(),
      conn: struct_like(Plug.Conn, path_info: ["send_200"]),
      plug: {__MODULE__, []}
    }

end

```

```

test "it should send `stop` events for normally completing requests", context do
  Req.get!(context.req, url: "/send_200")

  assert_receive [:telemetry, [:bandit, :request, :stop], measurements, metadata], 500

  assert measurements
    ~> %{
      resp_body_bytes: 0,
      duration: integer(max: System.convert_time_unit(1, :second, :native)),
      monotonic_time: integer(roughly: System.monotonic_time()),
      req_header_end_time: integer(roughly: System.monotonic_time()),
      resp_start_time: integer(roughly: System.monotonic_time()),
      resp_end_time: integer(roughly: System.monotonic_time())
    }

  assert metadata
    ~> %{
      connection_telemetry_span_context: reference(),
      telemetry_span_context: reference(),
      conn: struct_like(Plug.Conn, path_info: ["send_200"]),
      plug: {__MODULE__, []}
    }

end

```

```
test "it should send `stop` events for normally completing requests", context do
  Req.get!(context.req, url: "/send_200")

  assert_receive [:telemetry, [:bandit, :request, :stop], measurements, metadata], 500

  assert measurements
    ~> %{
      resp_body_bytes: 0,
      duration: integer(max: System.convert_time_unit(1, :second, :native)),
      monotonic_time: integer(roughly: System.monotonic_time()),
      req_header_end_time: integer(roughly: System.monotonic_time()),
      resp_start_time: integer(roughly: System.monotonic_time()),
      resp_end_time: integer(roughly: System.monotonic_time())
    }

  assert metadata
    ~> %{
      connection_telemetry_span_context: reference(),
      telemetry_span_context: reference(),
      conn: struct_like(Plug.Conn, path_info: ["send_200"]),
      plug: {__MODULE__, []}
    }

end
```

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable**

Tests should be **ubiquitous**



# Tests should be correct

A test is meaningless until you've seen it fail

- Red, green, refactor is basically gospel
  - The *very first* thing you should do is build a repro case
  - The *next* thing you should do is to codify that in a test
  - This is so important: **WATCH THAT TEST FAIL**
  - *Then* (and only then) can you get to fixing it
- This helps ensure that you're testing the right thing

# Tests should be correct

## Defence in depth

- More tests are a good thing
  - This *doesn't* mean 'blindly write more test cases'
- Scope your test plan to different levels of abstraction
  - Unit tests for important / subtle modules
  - Acceptance tests for overall behaviour (happy & common sad paths)
  - Mocks take on different roles in these cases



# Tests should be correct

## Mocks considered harmful

- Mocks intentionally diverge the system under test from production
  - You now have two problems
  - Perilously easy to gain false confidence
- If you *really* need to mock, Mox, Mimic & ex\_vcr are the way to go
  - Keep your mocks logic free. Input validation and static returns **only**
  - Explicitly test your mocks; they're real code too
- At least Mox mandates behaviours. Maybe this will get better with types?

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable**

Tests should be **ubiquitous**

# Tests should be **reliable**

Your tests should pass (or fail) 100% of the time

- Non-deterministic tests are a **huge** red flag
  - Is the non-determinism random (almost certainly not)
  - Is it due to isolation (suggests shared resources)
  - Is it due to load / timing (suggests a race condition)
- `--repeat-until-failure 10000` (1.17+) is handy

# Tests should be **reliable**

Isolate tests from one another

- Tests within a given module always run sequentially
- Run application anew every test (`start_supervised` in a `setup` block)
- Only one test in the module is ever running so `__MODULE__` is 'safe'
- Use `self()` & `assert_receive` to send messages to the test process
  - This also helps avoid race conditions against async code

```

defmodule PlugTest do
  test "it should send `stop` events for normally completing requests", context do
    TelemetryHelpers.attach_all_events(__MODULE__) |> on_exit()

    Req.get!(context.req, url: "/send_200")

    assert_receive {:telemetry, [:bandit, :request, :start], measurements, metadata}, 500
    assert measurements ~> %{monotonic_time: integer(roughly: System.monotonic_time()) }
  end
end

defmodule TelemetryHelpers do
  @events [...]

  def attach_all_events(plug) do
    ref = make_ref()
    :telemetry.attach_many(ref, @events, &__MODULE__.handle_event/4, {self(), plug})
    fn -> :telemetry.detach(ref) end
  end

  def handle_event(event, measurements, %{plug: {plug, _}} = metadata, {pid, plug}) do
    send(pid, {:telemetry, event, measurements, metadata})
  end
end

```

```

defmodule PlugTest do
  test "it should send `stop` events for normally completing requests", context do
    TelemetryHelpers.attach_all_events(__MODULE__) |> on_exit()

    Req.get!(context.req, url: "/send_200")

    assert_receive {:telemetry, [:bandit, :request, :start], measurements, metadata}, 500
    assert measurements ~> %{monotonic_time: integer(roughly: System.monotonic_time()) }
  end
end

defmodule TelemetryHelpers do
  @events [...]

  def attach_all_events(plug) do
    ref = make_ref()
    :telemetry.attach_many(ref, @events, &__MODULE__.handle_event/4, {self(), plug})
    fn -> :telemetry.detach(ref) end
  end

  def handle_event(event, measurements, %{plug: {plug, _}} = metadata, {pid, plug}) do
    send(pid, {:telemetry, event, measurements, metadata})
  end
end

```

```

defmodule PlugTest do
  test "it should send `stop` events for normally completing requests", context do
    TelemetryHelpers.attach_all_events(__MODULE__) |> on_exit()

    Req.get!(context.req, url: "/send_200")

    assert_receive {:telemetry, [:bandit, :request, :start], measurements, metadata}, 500
    assert measurements ~> %{monotonic_time: integer(roughly: System.monotonic_time()) }
  end
end

defmodule TelemetryHelpers do
  @events [...]

  def attach_all_events(plug) do
    ref = make_ref()
    :telemetry.attach_many(ref, @events, &__MODULE__.handle_event/4, {self(), plug})
    fn -> :telemetry.detach(ref) end
  end

  def handle_event(event, measurements, %{plug: {plug, _}} = metadata, {pid, plug}) do
    send(pid, {:telemetry, event, measurements, metadata})
  end
end

```



# Tests should be reliable

## Quiet & sane test output

- It's 8000x easier to with a test failure when you can see it in isolation
- Tests should be quiet with all output quelled or captured
- `@tag :capture_log` works wonders for this
- If you want to capture log output for tests, `import ExUnit.CaptureLog`
  - Sometimes unavoidable to need a `Process.sleep` in these cases
- Try not to overuse log capture; it can easily hide trouble in otherwise passing tests



Tests should be **clear**

Tests should be **correct**

Tests should be **reliable**

Tests should be **ubiquitous**

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable (& fast)**

Tests should be **ubiquitous**

# Tests should be reliable (& fast)

Run your tests in parallel

- `use ExUnit.Case async: true`
- Each test file will be run in parallel
- Tests within a single file are always run sequentially
- Overall run speed is limited by your slowest file
  - Prefer more, smaller test files

**Bandit's tests ~~should be~~ are fast**

**8x faster**

From ~45s to <6s

Tests should be **clear**

Tests should be **correct**

Tests should be **reliable (& fast)**

Tests should be **ubiquitous**

# Tests should be ubiquitous

You should be running CI

- GitHub Actions is someone else's computer, but free and amazing
  - Run tests + dialyzer + credo + others on every push to every branch
  - Matrix testing: test on all combinations of recent OTPs and Elixirs (or any other property)
- GitHub has a bunch of options to gate merges on CI

# Tests should be ubiquitous

Impossibly easy to get started

```
# Put this in .github/workflows/elixir.yml
```

```
name: Elixir CI
```

```
on:
```

```
  push:
```

```
    branches: [ main ]
```

```
  pull_request:
```

```
  workflow_dispatch:
```

```
jobs:
```

```
  test:
```

```
    uses: mtrudel/elixir-ci-actions/.github/workflows/test.yml@main
```

```
  lint:
```

```
    uses: mtrudel/elixir-ci-actions/.github/workflows/lint.yml@main
```

# Tests should be ubiquitous

## Impossibly easy to get started

The screenshot shows a GitHub Actions workflow run for the repository 'mtrudel / machete'. The workflow is named 'Elixir CI' and the specific run is 'Version bump to 0.3.8 #175'. The run was triggered via a push 2 weeks ago and is in a 'Success' status. The total duration of the run is 29 seconds. The workflow file is 'elixir.yml' and it runs on 'push' events. The workflow consists of two main jobs: 'test' and 'lint'. The 'test' job is expanded, showing a matrix of test runs for different Elixir versions and dependencies. The 'lint' job is also shown, with a single run for Elixir 1.18.x and dependency 27.x. The 'test' job is marked as completed, and the 'lint' job is also marked as completed.

github.com/mtrudel/machete/actions/runs/13061286185

mtrudel / machete

Type / to search

Code Issues Pull requests 1 Actions Security Insights Settings

← Elixir CI

✓ Version bump to 0.3.8 #175

Re-run all jobs

Summary

Jobs

- ✓ test
  - ✓ test (1.16.x, 25.x)
  - ✓ test (1.16.x, 26.x)
  - ✓ test (1.17.x, 25.x)
  - ✓ test (1.17.x, 26.x)
  - ✓ test (1.17.x, 27.x)
  - ✓ test (1.18.x, 25.x)
  - ✓ test (1.18.x, 26.x)
  - ✓ test (1.18.x, 27.x)
- ✓ lint
  - ✓ lint (1.18.x, 27.x)

Run details

elixir.yml

on: push

Matrix: lint / lint

- ✓ 1 job completed
- Show all jobs

Matrix: test / test

- ✓ test / test (1.16.x, 25.x) 16s
- ✓ test / test (1.16.x, 26.x) 11s
- ✓ test / test (1.17.x, 25.x) 13s



Tests that are **clear**

Tests that are **correct**

Tests that are **reliable (& fast)**

Tests that are **ubiquitous**

*fin*